

Weiwei Sun

 sunnweiwei@gmail.com  [@sunweiwei12](https://twitter.com/sunweiwei12)  [sunweiwei](https://www.linkedin.com/in/sunweiwei/)
 <https://sunnweiwei.github.io/>  [google scholar](#)



Overview

- Homepage**  <https://sunnweiwei.github.io/>
- Google Scholar**  <https://scholar.google.com/citations?user=hdUZbxgAAAAJ&hl=en>
- Research**  My recent research focuses on LLM agents for long-horizon tasks such as software engineering (SWE), algorithm development and deep research. I have trained agents that actively manage context [1] and collaborate with users in a proactive and personalized manner [2]. I have designed agents for automated machine learning engineering and algorithm search [3], [4]. I have also worked on LLM behavior attribution and transparency [5], as well as retrieval-augmented generation (RAG) and generative search [6], [7].

Education

- 2024.8 – now  **Carnegie Mellon University**
PhD in Language Technologies Institute
Supervisor: Prof. Yiming Yang
- 2017.9 – 2023.12  **Shandong University**
B.S. and M.S. in Computer Science and Technology
Supervisor: Prof. Zhaochun Ren

Research Experience

- 2024.8 – now  **Carnegie Mellon University.** Pittsburgh.
Research Assistant | Advisor: Yiming Yang.
Research: LLM Agent, AI4Science.
- 2025.5 – 2025.8  **ByteDance Seed.** San Jose.
Research Intern | Mentor: Jiecao Chen.
Research: LLM Agent, RL.
- 2024.7 – 2024.8  **Vector Institute.** Toronto.
Research Intern | Mentor: Colin Raffel.
Research: Training Data Attribution.
- 2024.3 – 2024.7  **University of Amsterdam.** Amsterdam.
Research Assistant | Mentor: Maarten de Rijke.
Research: Generative Retrieval.
- 2022.9 – 2023.12  **Baidu Inc.** Beijing.
Research Intern | Mentor: Dawei Yin.
Research: LLMs for Search.
- 2021.3 – 2021.9  **JD.com.** Beijing.
Research Intern | Mentor: Hongshen Chen.
Research: Question Answering.
- 2019.9 – 2021.3  **Shandong University.** Qingdao.
Research Assistant | Mentor: Zhaochun Ren.
Research: RAG, LLM Evaluation.

- 1 **Weiwei Sun**, M. Lu, Z. Ling, K. Liu, X. Yao, Y. Yang, and J. Chen, “Scaling long-horizon llm agent via context-folding,” in *arXiv*, 2025. ⚡ URL: <https://arxiv.org/abs/2510.11967>.
- 2 **Weiwei Sun**, X. Zhou, W. Du, X. Wang, S. Welleck, G. Neubig, M. Sap, and Y. Yang, “Training proactive and personalized llm agents,” in *arXiv*, 2025. ⚡ URL: <https://arxiv.org/abs/2511.02208>.
- 3 S. Li*, **Weiwei Sun***, S. Li, A. Talwalkar, and Y. Yang, “Toward community-driven agents for machine learning engineering,” in *MTI-LLM@NeurIPS*, 2025. ⚡ URL: <https://arxiv.org/abs/2506.20640>.
- 4 **Weiwei Sun**, S. Feng, S. Li, and Y. Yang, “Co-bench: Benchmarking language model agents in algorithm search for combinatorial optimization,” in *AAAI*, 2026. ⚡ URL: <https://arxiv.org/abs/2504.04310>.
- 5 **Weiwei Sun**, H. Liu, N. Kandpal, C. Raffel, and Y. Yang, “Enhancing training data attribution with representational optimization,” in *NeurIPS*, 2025, *Spotlight*. ⚡ URL: <https://arxiv.org/abs/2505.18513>.
- 6 Y. Chen, L. Yan, **Weiwei Sun**, X. Ma, Y. Zhang, S. Wang, D. Yin, Y. Yang, and J. Mao, “Improving retrieval-augmented generation through multi-agent reinforcement learning,” in *NeurIPS*, 2025. ⚡ URL: <https://arxiv.org/abs/2501.15228>.
- 7 **Weiwei Sun**, K. Kong, X. Ma, S. Wang, D. Yin, M. de Rijke, Z. Ren, and Y. Yang, “Zerogr: A generalizable and scalable framework for zero-shot generative retrieval,” in *arXiv*, 2025. ⚡ URL: <https://arxiv.org/abs/2510.10419>.
- 8 S. Feng*, **Weiwei Sun***, S. Li, A. Talwalkar, and Y. Yang, “Frontierco: A comprehensive evaluation of contemporary machine learning solvers for combinatorial optimization,” in *AI4MATH@ICML*, 2025. ⚡ URL: <https://arxiv.org/abs/2505.16952>.
- 9 **Weiwei Sun**, Z. Shi, J. Wu, L. Yan, X. Ma, Y. Liu, M. Cao, D. Yin, and Z. Ren, “Mair: A massive benchmark for evaluating instructed retrieval,” in *EMNLP*, 2024. ⚡ URL: <https://arxiv.org/abs/2410.10127>.
- 10 **Weiwei Sun**, Z. Chen, X. Ma, L. Yan, S. Wang, P. Ren, Z. Chen, D. Yin, and Z. Ren, “Instruction distillation makes large language models efficient rankers,” in *CCIR and GenRec@CIKM*, 2023. ⚡ URL: <https://arxiv.org/abs/2311.01555>.
- 11 **Weiwei Sun**, L. Yan, Z. Chen, S. Wang, H. Zhu, P. Ren, Z. Chen, D. Yin, M. de Rijke, and Z. Ren, “Learning to tokenize for generative retrieval,” in *NeurIPS*, 2023. ⚡ URL: <https://openreview.net/pdf?id=UKd6dpVGdu>.
- 12 **Weiwei Sun**, L. Yan, X. Ma, S. Wang, P. Ren, Z. Chen, D. Yin, and Z. Ren, “Is chatgpt good at search? Investigating large language models as re-ranking agents,” in *EMNLP*, 2023, *Outstanding paper award*. ⚡ URL: <https://arxiv.org/abs/2304.09542>.
- 13 **Weiwei Sun**, H. Cai, H. Chen, P. Ren, Z. Chen, M. de Rijke, and Z. Ren, “Answering ambiguous questions via iterative prompting,” in *ACL*, 2023, pp. 7669–7683. ⚡ URL: <https://arxiv.org/abs/2307.03897>.
- 14 **Weiwei Sun**, P. Ren, and Z. Ren, “Generative knowledge selection for knowledge-grounded dialogues,” in *Findings of EACL*, 2023, pp. 2077–2088. ⚡ URL: <https://arxiv.org/abs/2304.04836>.
- 15 **Weiwei Sun**, Z. Shi, S. Gao, P. Ren, M. de Rijke, and Z. Ren, “Contrastive learning reduces hallucination in conversations,” in *AAAI*, 2023, pp. 13 618–13 626. ⚡ URL: <https://arxiv.org/abs/2212.10400>.
- 16 **Weiwei Sun**, S. Guo, S. Zhang, P. Ren, Z. Chen, M. de Rijke, and Z. Ren, “Metaphorical user simulators for evaluating task-oriented dialogue systems,” *ACM Transactions on Information Systems*, vol. 42, 2023. ⚡ URL: <https://arxiv.org/abs/2204.00763>.

- 17 Weiwei Sun, S. Zhang, K. Balog, Z. Ren, P. Ren, Z. Chen, and M. de Rijke, “Simulating user satisfaction for the evaluation of task-oriented dialogue systems,” in *SIGIR*, 2021, pp. 2499–2506. ⚒ URL: <https://arxiv.org/abs/2105.03748>.
- 18 Weiwei Sun, C. Meng, Q. Meng, Z. Ren, P. Ren, Z. Chen, and M. de Rijke, “Conversations powered by cross-lingual knowledge,” in *SIGIR*, 2021, pp. 1442–1451. ⚒ URL: <https://irlab.science.uva.nl/wp-content/papercite-data/pdf/sun-2021-conversations.pdf>.
- 19 M. Lu, Weiwei Sun, W. Du, Z. Ling, X. Yao, K. Liu, and J. Chen, “Scaling llm multi-turn rl with end-to-end summarization-based context management,” in *arXiv*, 2025. ⚒ URL: <https://arxiv.org/abs/2510.06727>.
- 20 W. Liu, X. Ma, Weiwei Sun, Y. Zhu, Y. Li, D. Yin, and Z. Dou, “Reasonrank: Empowering passage ranking with strong reasoning ability,” in *arXiv*, 2025. ⚒ URL: <https://arxiv.org/abs/2508.07050>.
- 21 S. Li, T. Marwah, J. Shen, Weiwei Sun, A. Risteski, Y. Yang, and A. Talwalkar, “Codepde: Benchmarking llms’ abilities to solve pdes through code generation,” in *arXiv*, 2025. ⚒ URL: <https://arxiv.org/abs/2505.08783>.
- 22 Z. Shi, L. Yan, Weiwei Sun, Y. Feng, P. Ren, X. Ma, S. Wang, D. Yin, M. de Rijke, and Z. Ren, “Direct retrieval-augmented optimization: Synergizing knowledge selection and language models,” in *arXiv*, 2025. ⚒ URL: <https://arxiv.org/abs/2505.03075>.
- 23 Z. Zhang, X. Ma, Weiwei Sun, P. Ren, Z. Chen, S. Wang, D. Yin, M. de Rijke, and Z. Ren, “Replication and exploration of generative retrieval over dynamic corpora,” in *SIGIR*, 2025.
- 24 Q. Liu, D. Haozhe, Y. Chen, Q. Lu, Weiwei Sun, and J. Mao, “Llm4ranking: An easy-to-use framework of utilizing large language models for document reranking,” in *arXiv*, 2025.
- 25 Y. Chen, Q. Liu, Y. Zhang, Weiwei Sun, D. Shi, J. Mao, and D. Yin, “Tourrank: Utilizing large language models for documents ranking with a tournament-inspired strategy,” in *WWW*, 2025. ⚒ URL: <https://arxiv.org/abs/2406.11678>.
- 26 J. Hao, Weiwei Sun, X. Xin, Q. Meng, Z. Chen, P. Ren, and Z. Ren, “Meft: Memory-efficient fine-tuning through sparse adapter,” in *ACL*, 2024. ⚒ URL: <https://arxiv.org/abs/2406.04984>.
- 27 Z. Shi, S. Zhang, Weiwei Sun, S. Gao, P. Ren, and Z. C. abd Zhaochun Ren, “Generate-then-ground in retrieval-augmented generation for multi-hop question answering,” in *ACL*, 2024. ⚒ URL: <https://arxiv.org/abs/2406.14891>.
- 28 Y. Wang, Z. Ren, Weiwei Sun, J. Yang, Z. Liang, X. Chen, R. Xie, S. Yan, X. Zhang, P. Ren, Z. Chen, and X. Xin, “Content-based collaborative generation for recommender systems,” in *CIKM*, 2024. ⚒ URL: <https://arxiv.org/abs/2403.18480>.
- 29 T. Ju, Weiwei Sun, W. Du, X. Yuan, Z. Ren, and G. Liu, “How large language models encode context knowledge? A layer-wise probing study,” in *LREC-Coling*, 2024. ⚒ URL: <https://arxiv.org/abs/2402.16061>.
- 30 Y. Zhao, L. Yan, Weiwei Sun, G. Xing, S. Wang, C. Meng, Z. Cheng, Z. Ren, and D. Yin, “Improving the robustness of large language models via consistency alignment,” in *LREC-Coling*, 2024. ⚒ URL: <https://arxiv.org/abs/2403.14221>.
- 31 Y. Zhao, L. Yan, Weiwei Sun, G. Xing, C. Meng, S. Wang, Z. Cheng, Z. Ren, and D. Yin, “Knowing what llms do not know: A simple yet effective self-detection method,” *NAACL*, 2024. ⚒ URL: <https://arxiv.org/abs/2310.17918>.
- 32 Y. Zhao, L. Yan, Weiwei Sun, C. Meng, S. Wang, Z. Cheng, Z. Ren, and D. Yin, “Diqad: A benchmark dataset for end-to-end open-domain dialogue assessment,” in *Findings of EMNLP*, 2023. ⚒ URL: <https://aclanthology.org/2023.findings-emnlp.1010.pdf>.

- 33 Z. Shi, **Weiwei Sun**, S. Zhang, Z. Zhang, P. Ren, and Z. Ren, “Rade: Reference-assisted dialogue evaluation for open-domain dialogue,” in *ACL*, 2023, pp. 12 856–12 875. ⚡ URL: <https://arxiv.org/abs/2309.08156>.
- 34 S. Guo, S. Zhang, **Weiwei Sun**, P. Ren, Z. Chen, and Z. Ren, “Towards explainable conversational recommender systems,” in *SIGIR*, 2023, pp. 2786–2795. ⚡ URL: <https://arxiv.org/abs/2305.18363>.
- 35 C. Meng, P. Ren, Z. Chen, **Weiwei Sun**, Z. Ren, Z. Tu, and M. de Rijke, “Dukenet: A dual knowledge interaction network for knowledge-grounded conversation,” in *SIGIR*, 2020, pp. 1151–1160. ⚡ URL: <https://ir.sdu.edu.cn/~zhaochunren/papers/meng-2020-dukenet.pdf>.

Awards and Achievements

- 2025 📚 **LTI Presidential Fellowship**, CMU.
- 2024 📚 **Baidu Scholarship**, Global Top-10, Baidu Inc.
- 2023 📚 **Outstanding Paper Award**, EMNLP 2023.
- 📘 **Presidential Scholarship**, Top 0.1%, Shandong University.
- 📘 **National Scholarship**, Top 0.2%, Shandong University.
- 2021 📚 **Outstanding Graduate with Honorary Degree**, Shandong University.
- 2019 📚 **National Second Price**, Mathematical Contest in Modeling (CUMCM).
- 2017 – 2023 📚 **College Academic Scholarship**, Shandong University.

Selected Projects

- 2025 📚 **PPP-Agent**: Agent optimized for productive, proactive, and personalized agent–user interactions. The PPP RL framework outperforms GPT-5 on vague user prompts.
<https://arxiv.org/abs/2511.02208> <https://github.com/sunnweiwei/PPP-Agent>
- 📘 **Context-Folding**: Agent that actively compresses and folds its context. RL algorithm enables end-to-end optimization with significant improv. on SWE-Bench and BrowseComp.
<https://arxiv.org/abs/2510.11967> <https://github.com/sunnweiwei/FoldAgent>
- 📘 **CoMind**: Community-driven multi-agent system for automated machine learning engineering. SOTA results on MLE-Bench and rank top 7.4% of ongoing Kaggle competitions.
<https://arxiv.org/abs/2506.20640> <https://github.com/comind-ml/CoMind>
- 2023 📚 **RankGPT**: First work to use LLMs for zero-shot listwise text ranking. SOTA results on major IR benchmarks and is adopted by LangChain / LlamaIndex / etc.
<https://arxiv.org/abs/2304.09542> <https://github.com/sunnweiwei/RankGPT>

Academic Service

- | | |
|-----------|---|
| Reviewer | 📘 EMNLP 2022, ACL 2023, SIGIR 2023, SIGIR-AP 2023, ECML/PKDD 2023, IPM 2023, WSDM 2023, EMNLP 2023, TALLIP 2023, ACM TOIS 2024, SIGIR 2024, ECML/PKDD 2024, EMNLP 2024 (SAC), ACL ARR 2024 June SAC, CIKM 2024, NeurIPS 2024, WSDM 2025, ICLR 2025, AISTATS 2025, ICML 2025, NeurIPS 2025 (Top Reviewer), AAAI 2026, ICLR 2026. |
| Volunteer | 📘 YSSNLP 2022, NLPCC 2022, SIGIR-AP 2023. |
| Tutorial | 📘 WWW 2024, Recent Advances in Generative Information Retrieval. |
| Workshop | 📘 The 2025 LTI Student Research Symposium. Link. |